CIFAR-10 DiT · denoising learning curves

45,000 train / 5,000 validation · fixed validation noise and timestep seed 2027 · same split in every run

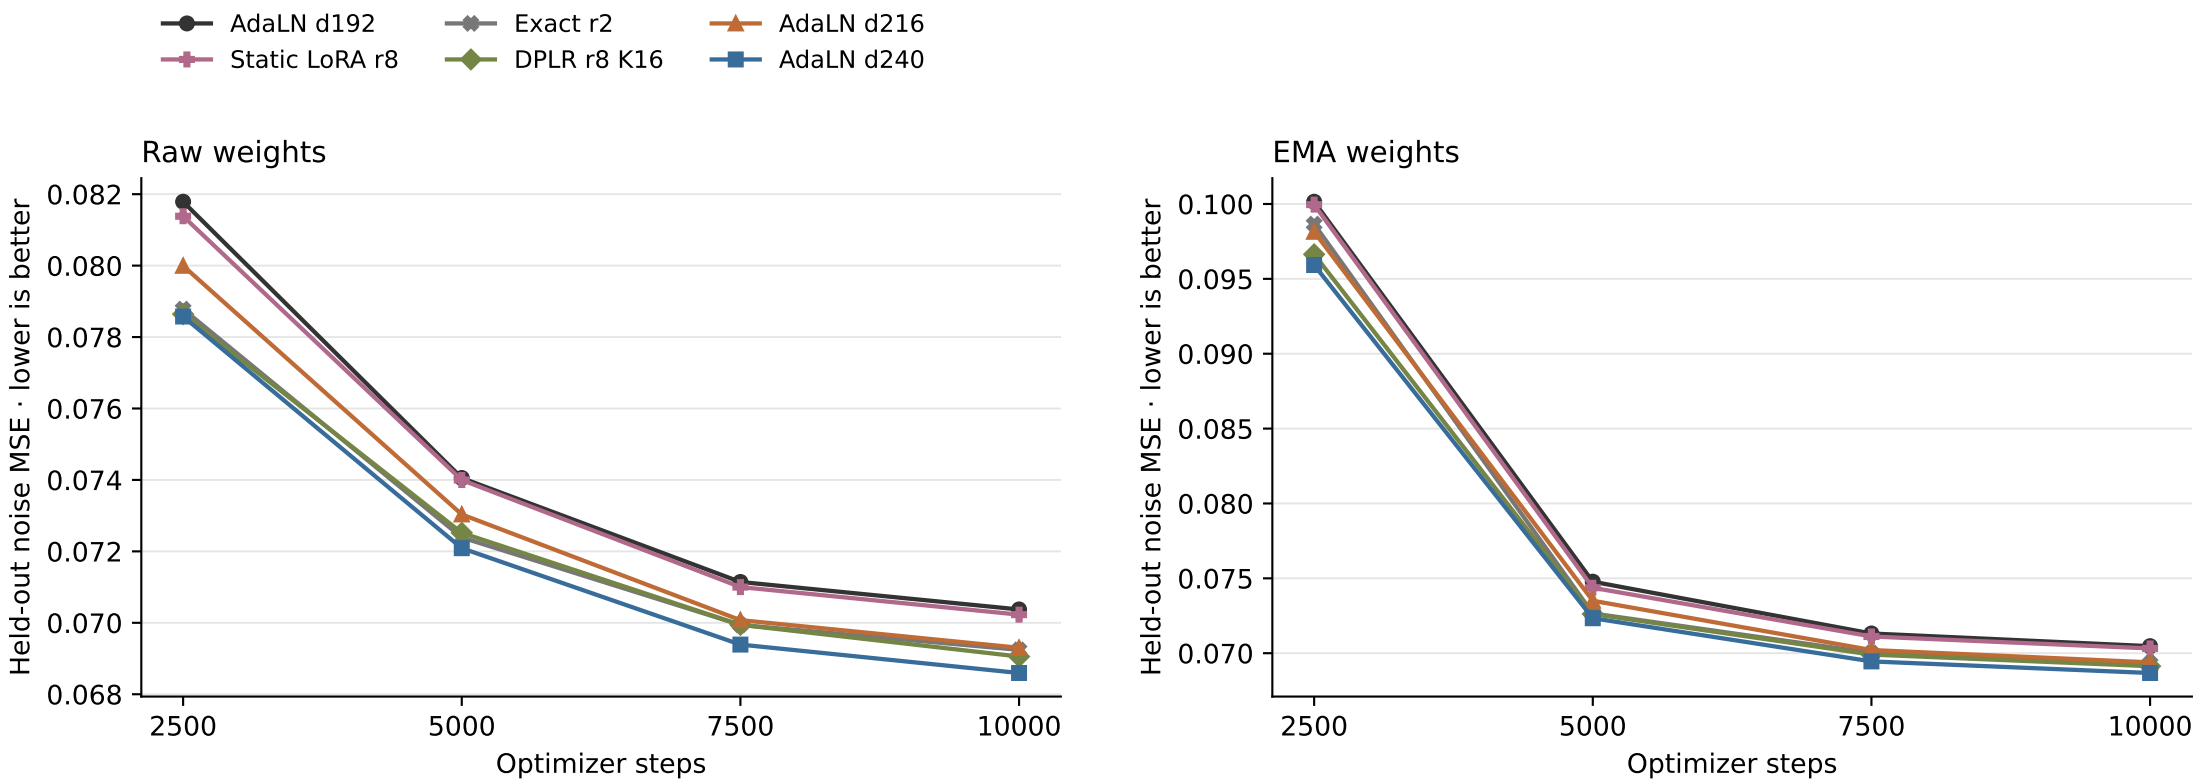

Initial noise MSE is approximately 1 for every zero-output model; step 0 omitted to show the learning region.

Single seed and common untuned diffusion recipe. Exact adapters use LR multiplier 0.1 transferred from ViT diagnostics.

Denoising loss and global FID do not measure generated class accuracy or sequence-conditioned attention.